

Internet of Things Metadata

Aslak Johansen asjo@mmmi.sdu.dk


Feb 8, 2021

Part 0: Introduction

Motivation: A Cyber-Physical System

Consider a system consisting of

- ▶ A building with temperature sensors, actuators and setpoints per room
- ▶ A dashboard application showing a floormap colorcoded according to measured temperature
- ▶ A thermostat application which controls the temperature of each room



Motivation: Connecting the Data and Control Planes

Each point (sensor, actuator and setpoint) is named through a UUID but have different access rights.

How can we map the two applications to the points of a building?

1. Scatter constants throughout the code of each application
2. Introduce a per-application data structure:
room name \mapsto (sensor|actuator|setpoint) \mapsto UUID
3. Share that structure

What happens when we add another building?

What happens when we add another application concerned with the power consumption of the lighting of a single floor of the building?

Problem

We want to write applications for buildings.


Attractive qualities:

- ▶ ***Portability*** An application can be executed on a many buildings without modification.
- ▶ ***Maintainability*** A change in a building does not translate to a need for changing an application.

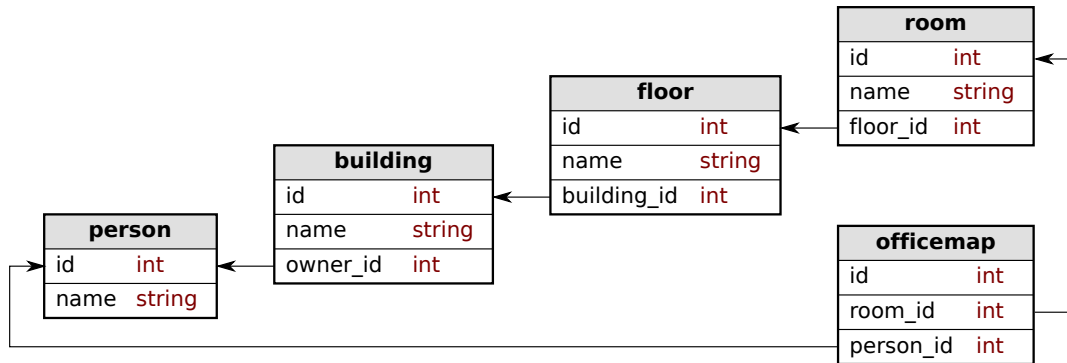
How do we accomplish this?

Approach: A Narrow Waist


The labor intensive job of mapping out the equipment, data streams and relations between those may be shared between a portfolio of applications




Relational Database



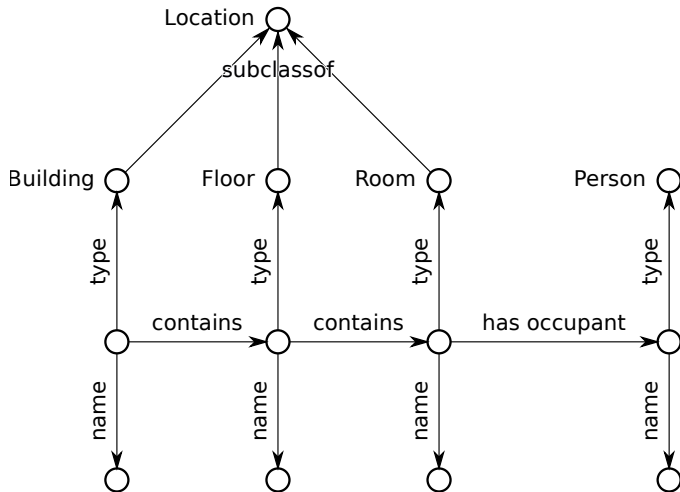
Object Database



Graph Database



Ontology-Based Information Model



Part 1: Semantics

Context: Semantics

"cec78d2a-14b8-41d6-bf46-c6a4d07574a7 is in room U168"

What is the *intension* of this statement?

- ▶ *cec78d2a-14b8-41d6-bf46-c6a4d07574a7* is a UUID
- ▶ *cec78d2a-14b8-41d6-bf46-c6a4d07574a7* identifies some data stream
- ▶ *U168* identifies a room
- ▶ The data stream identified by *cec78d2a-14b8-41d6-bf46-c6a4d07574a7* originates from some equipment in the room identified by *U168*

One of these intensions is expressed, the others are merely implied

Ontologies

According to Wikipedia:

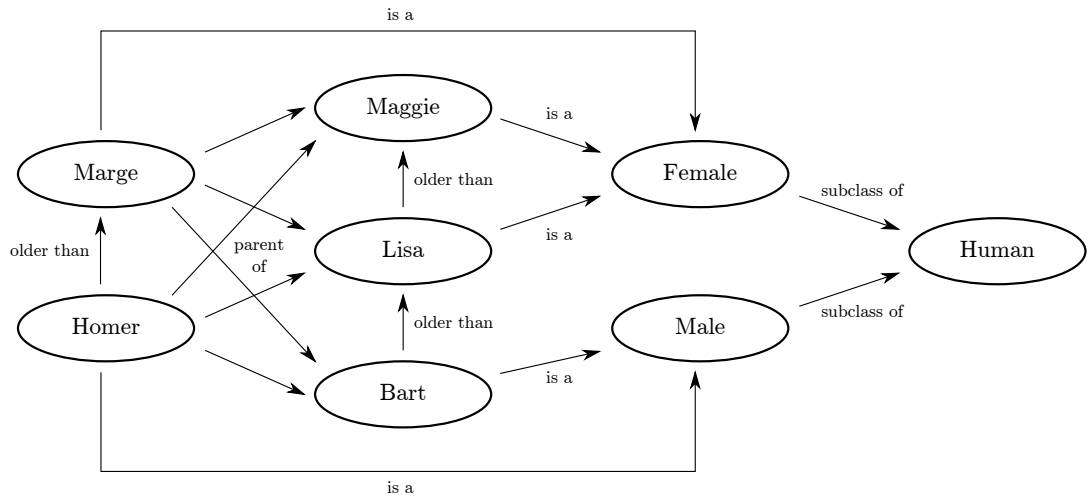
*"In computer science and information science, an ontology is a **formal naming and definition of the types, properties, and interrelationships of the entities that really exist in a particular domain of discourse.**"*

Lets look into what this means :

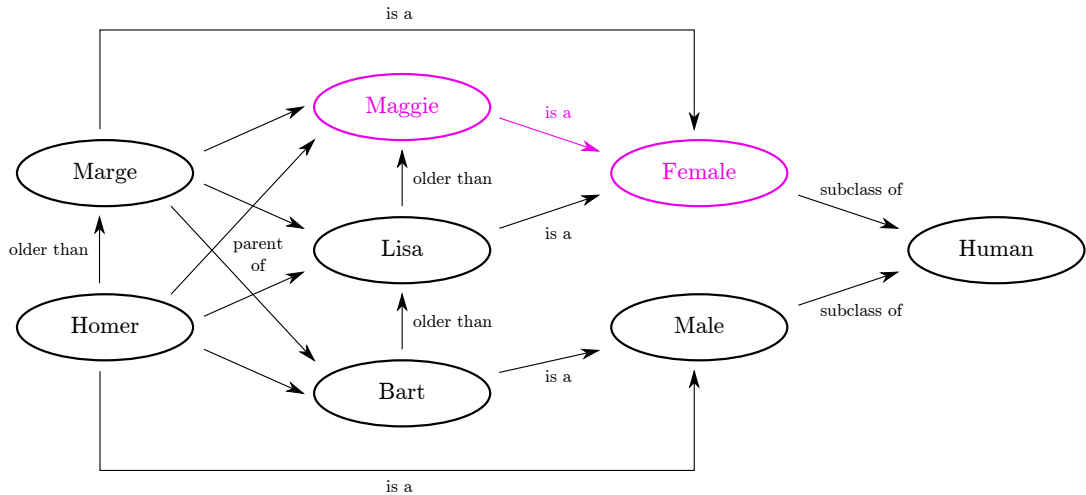
- ▶ **Formal Naming and Definition** A formalized language is employed
- ▶ **Types, Properties and Interrelationships of Entities** Things are defined through their type, parameters and relationships to other things
- ▶ **Entities Existing in a Particular Domain** Ontologies have limited coverage and focus on a particular field

Part 2: The Resource Description Framework

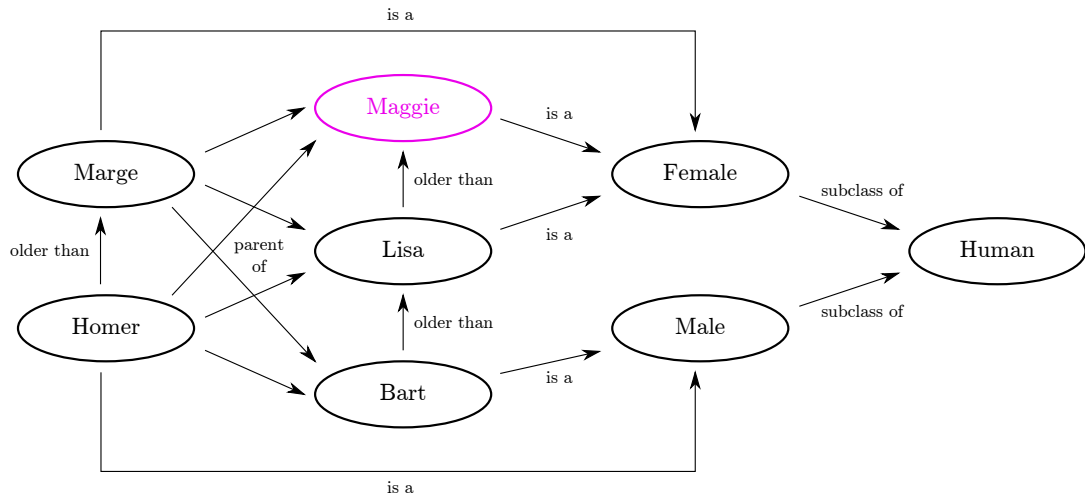
RDF Example



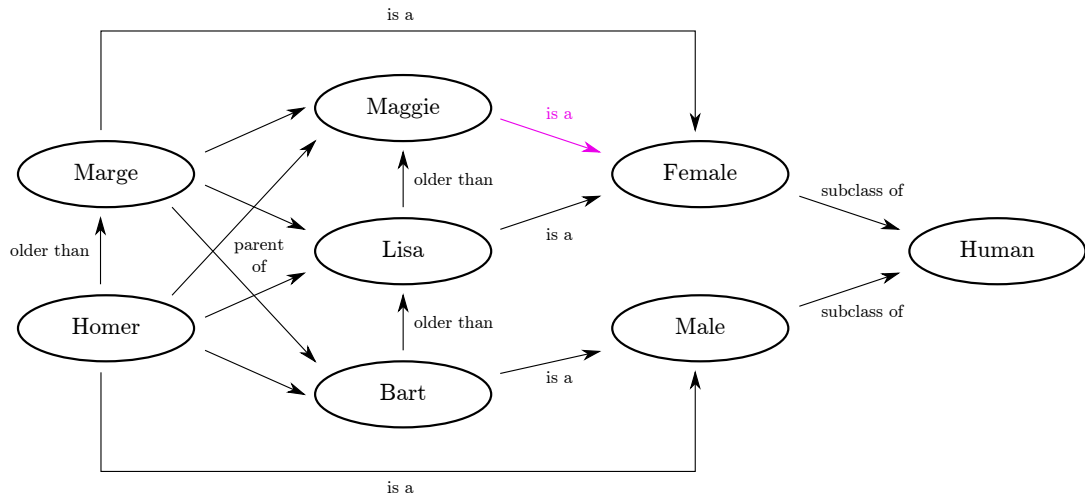
RDF Example



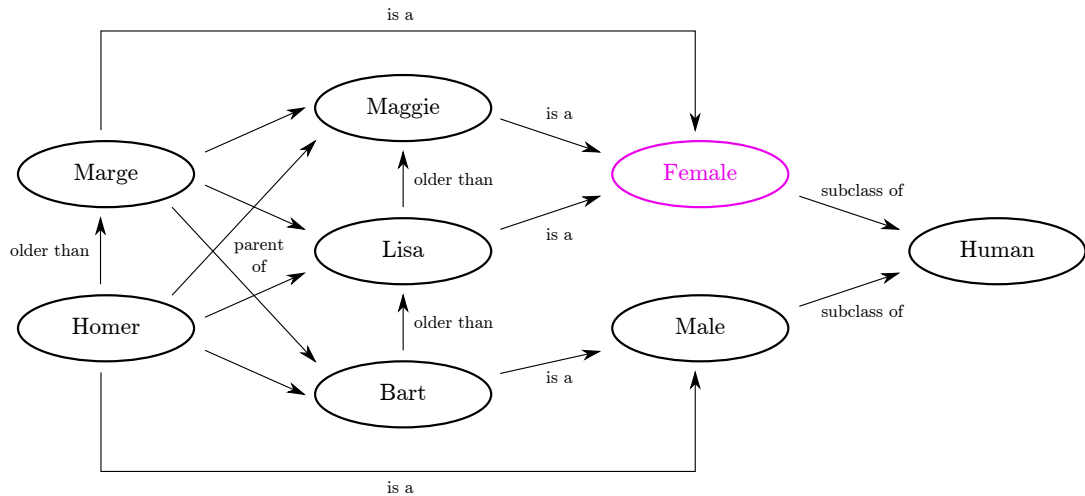
RDF Example




RDF Example



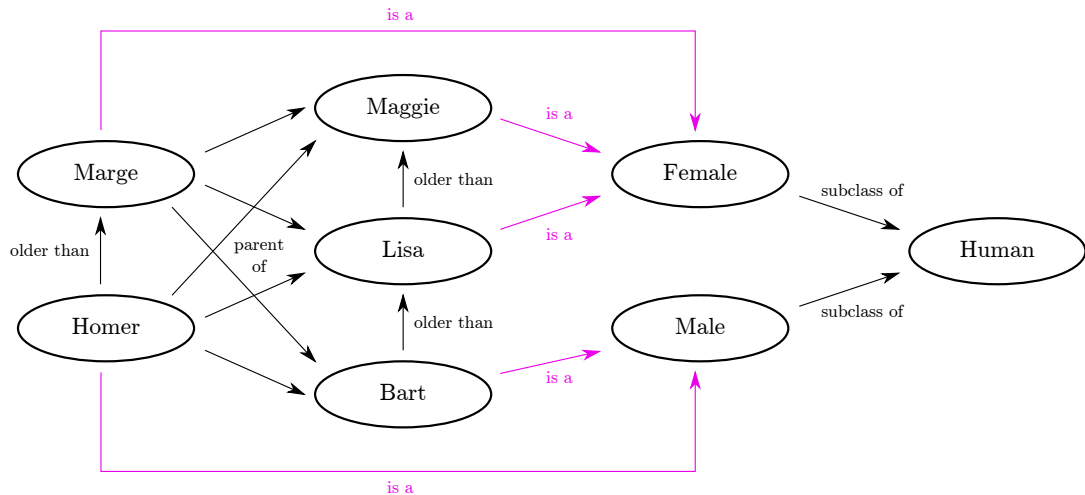
RDF Example



RDF Example



RDF Example



RDF - The Resource Description Framework

An ontology defined through triples

- ▶ Triples: subject \times predicate \times object
- ▶ A triple store is a model
- ▶ Namespaces
- ▶ Subclasses and subproperties

Wherein lies the information?

Hint: It is not in the name

<i>Subject</i>	<i>Predicate</i>	<i>Object</i>
Marge	parent of	Bart
Marge	is a	Female
Bart	is a	Male
Female	subclass of	Human
Male	subclass of	Human

In the relationships: A triple is a fact!

Turtle - The Terse RDF Triple Language

One (out of many) standardized ways of marshalling RDF.

File extension: .ttl

```
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix brick: <https://brickschema.org/schema/1.1.0/Brick#> .
@prefix demo: <http://ontologies.sdu.dk/Demo#> .

demo:floor a brick:Floor .

demo:room a brick:Room ;
    rdfs:label "U168" ;
    brick:isLocatedIn demo:floor .
```

OWL: The Web \leftrightarrow Ontology Language

A set of extensions to RDF

Adds formal semantics:

- ▶ **Constraints** e.g., on cardinality
- ▶ **Relationship Inferences** transitive, symmetric, inverse ...
- ▶ **Equivalence Testing** whether two concepts are similar enough
- ▶ **Subsumption Testing** whether one concept is more general than another
- ▶ ...

Closed world assumption: What is not known to be true must be false!

Part 3:

Modeling Buildings with Brick

Brick An Ontology for Building Metadata

Builds on OWL

Subject	Predicate	Object
my:building	type	brick:Building
my:building/floor2	type	brick:Floor
my:building/room1	type	brick:Room
my:building	contains	my:building/floor2
my:building/floor2	contains	my:building/room1
my:building/room1	label	Literal("U168")

Quizztime: What is the name of the building?

Brick Concepts

Definitions for

- ▶ Typing (Room, Temperature Sensor)
- ▶ Flows (feeds, feeds air)
- ▶ Physical encapsulation (contains)
- ▶ Control (controls)

Uncertainty

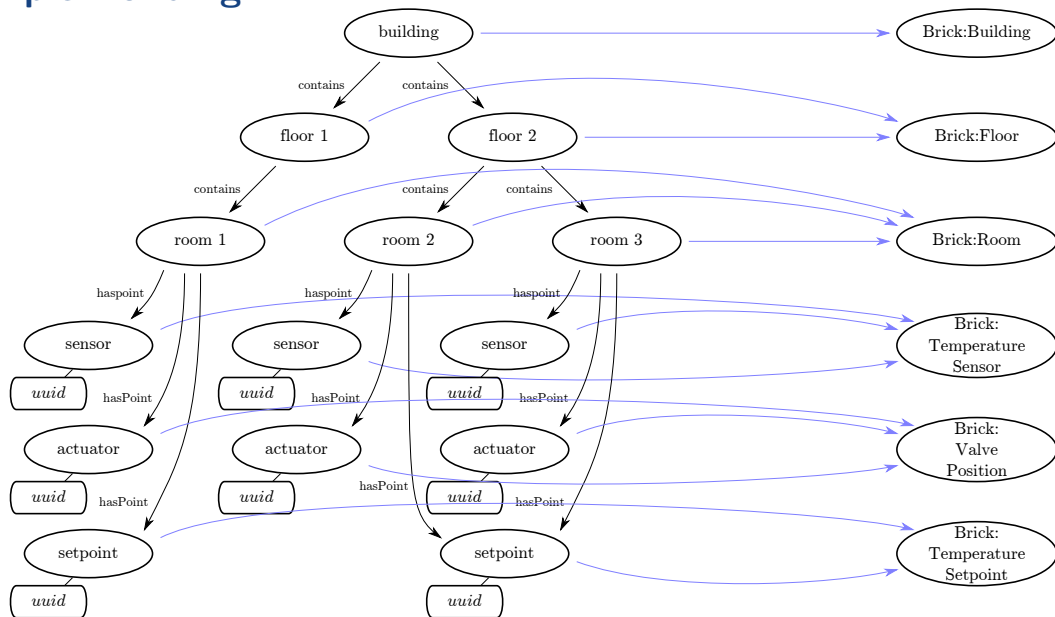
- ▶ Hierarchy of types (sensor < temperature sensor < celcius temperature sensor)

May be considered a graph ... but isn't

Quizz:

- ▶ What is a graph?
- ▶ How does a tripplestore differ?

Sample Building



Part 4:

Model Construction

Demo: Add Model

```
#!/usr/bin/env python3
```

```
from wrapping import *
```

```
g = model()
```

All code available at <https://github.com/aslakjohansen/sdu-iot-brick>

Demo: Add Building

```
building = N['/building']  
g.add((building, RDF.type, BRICK['Building']))
```

Demo: Add Floors

```
floors = []  
for floor_number in range(2):  
    floor = N['building/floors/%u' % floor_number]  
    g.add((floor, RDF.type, BRICK['Floor']))  
    g.add((building, BRICK.contains, floor))  
    floors.append(floor)
```


Demo: Add Room Mapping

```
rooms = {  
    'room 1': {  
        'floor': 0,  
        'temp-sensor': 'd0fcec33-af08-44a8-b74b-7a51f1902d13',  
        'temp-actuator': '5c2e2bdb-5142-464d-ad91-ae483633dfd6',  
        'temp-setpoint': 'c087d1ba-03f6-4c60-85de-b95d82e63248',  
    },  
    'room 2': {  
        'floor': 1,  
        'temp-sensor': '18407632-8e81-4bb9-ae77-4ecc96f30f46',  
        'temp-actuator': 'ec883cad-3235-46d5-a901-7bae169dda0a',  
        'temp-setpoint': 'd2b54605-58b1-44ef-89f9-0a9adabbfbb1',  
    },  
    'room 3': {  
        'floor': 1,  
        'temp-sensor': 'e528f733-2e3a-4f65-93aa-6fb2e36d4b27',  
        'temp-actuator': '9be52aed-0e12-49b5-8680-47778b9b2adf',  
        'temp-setpoint': 'd2b54605-58b1-44ef-89f9-0a9adabbfbb1',  
    },  
}
```

Demo: Add Rooms

```
roommap = {}  
for roomname in rooms:  
    data = rooms[roomname]  
    room = N['building/rooms/%s' % roomname.replace(' ', '_')]  
    g.add((room, RDF.type, BRICK['Room']))  
    g.add((room, BRICK.label, Literal(roomname)))  
    g.add((floors[data['floor']], BRICK.contains, room))  
    roommap[roomname] = room
```

Demo: Add Points

```
for roomname in rooms:
    data = rooms[roomname]
    room = roommap[roomname]

    sensor = N['building/rooms/%s/temp-sensor' % roomname.replace(' ', '_')]
    g.add((sensor, RDF.type, BRICK['Temperature_Sensor']))
    g.add((sensor, BRICK.label, Literal(data['temp-sensor'])))
    g.add((sensor, BRICK.pointOf, room))


    setpoint = N['building/rooms/%s/temp-setpoint' % roomname.replace(' ', '_')]
    g.add((setpoint, RDF.type, BRICK['Temperature_Setpoint']))
    g.add((setpoint, BRICK.label, Literal(data['temp-setpoint'])))
    g.add((setpoint, BRICK.pointOf, room))

    actuator = N['building/rooms/%s/temp-actuator' % roomname.replace(' ', '_')]
    g.add((actuator, RDF.type, BRICK['Radiator_Valve_Position']))
    g.add((actuator, BRICK.label, Literal(data['temp-actuator'])))
    g.add((actuator, BRICK.pointOf, room))
```

Part 5:

Querying a Model

Pattern Matching: SparQL Query for Dashboard Application



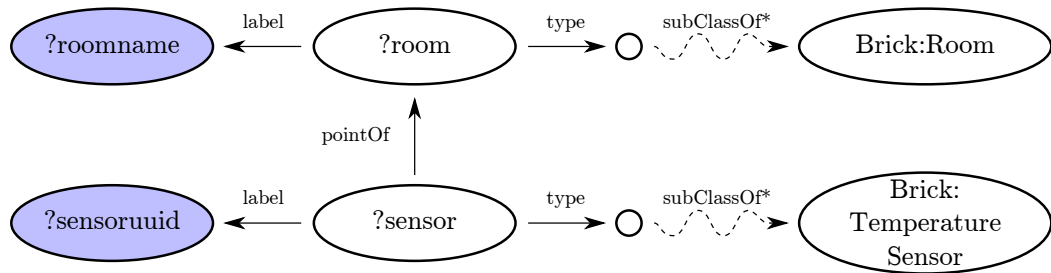
Dashboard Query SparQL Form

```
SELECT DISTINCT ?room_name ?sensor_uuid
WHERE {
    ?room      rdf:type/brick:subClassOf* brick:Room .
    ?sensor    rdf:type/brick:subClassOf* brick:Temperature_Sensor .

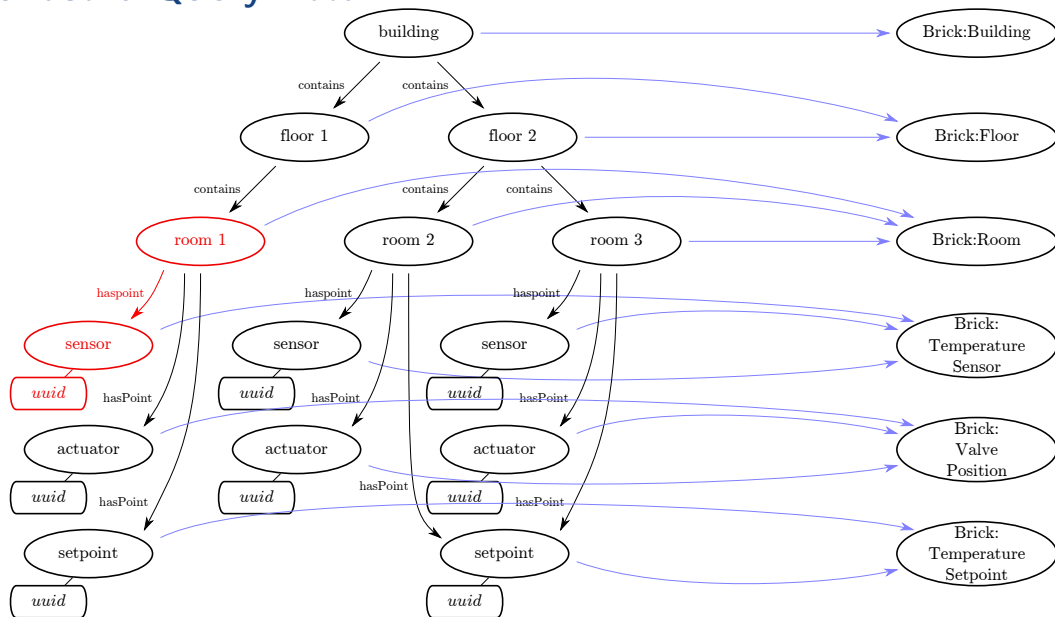
    ?sensor    brick:pointOf ?room .

    ?room      brick:label ?room_name .
    ?sensor    brick:label ?sensor_uuid .
}
```

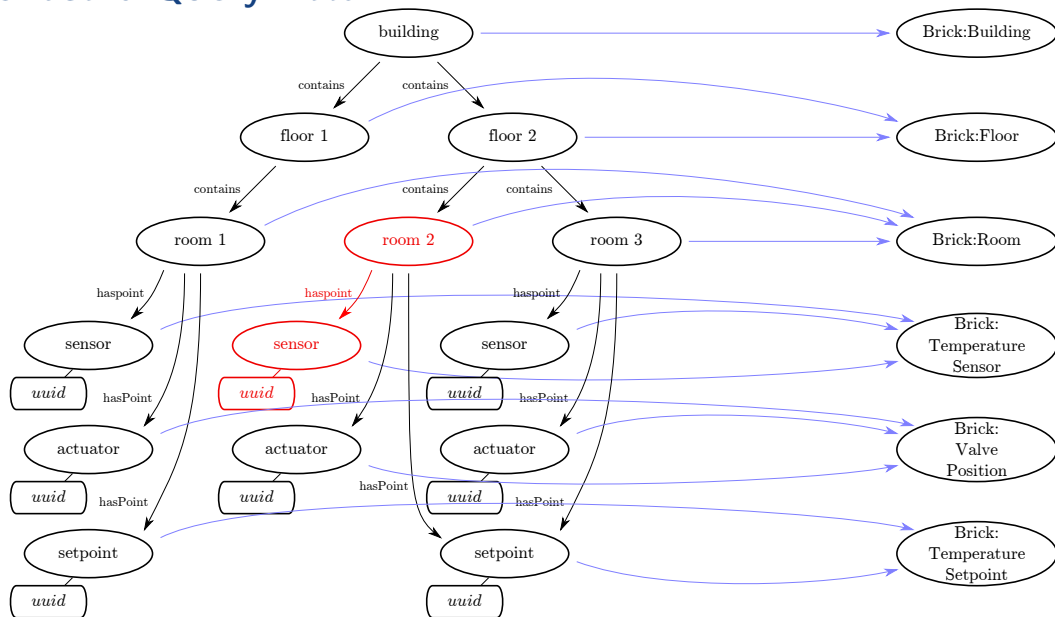
Dashboard Query Visual Form



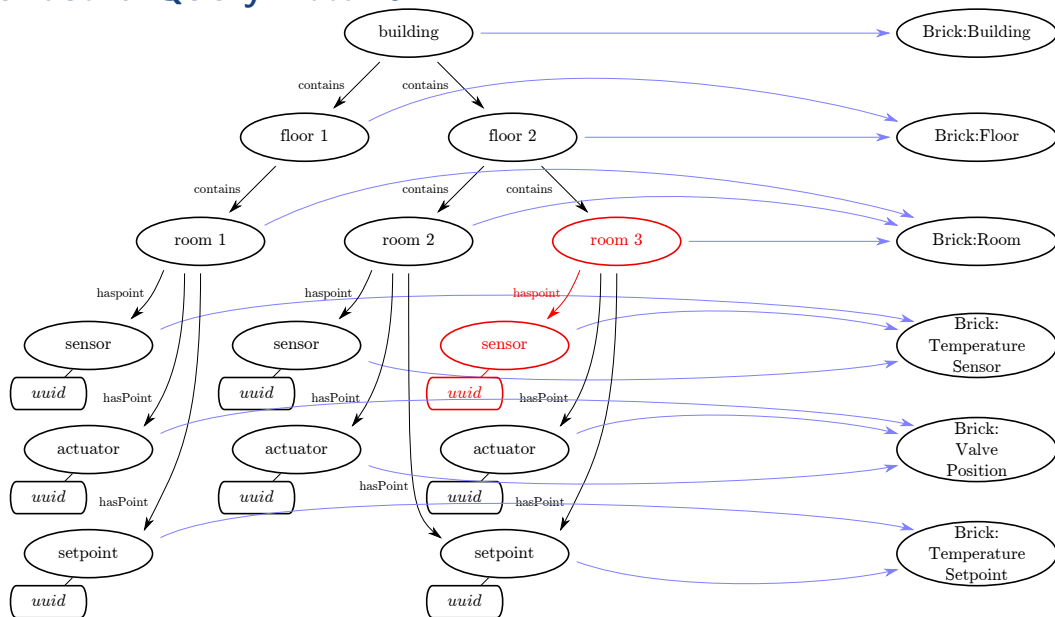
Dashboard Query Match 1



Dashboard Query Match 2



Dashboard Query Match 3



Dashboard Query Result Set

```
[  
  [  
    "room 1",  
    "d0fcec33-af08-44a8-b74b-7a51f1902d13"  
  ],  
  [  
    "room 2",  
    "18407632-8e81-4bb9-ae77-4ecc96f30f46"  
  ],  
  [  
    "room 3",  
    "e528f733-2e3a-4f65-93aa-6fb2e36d4b27"  
  ]  
]
```

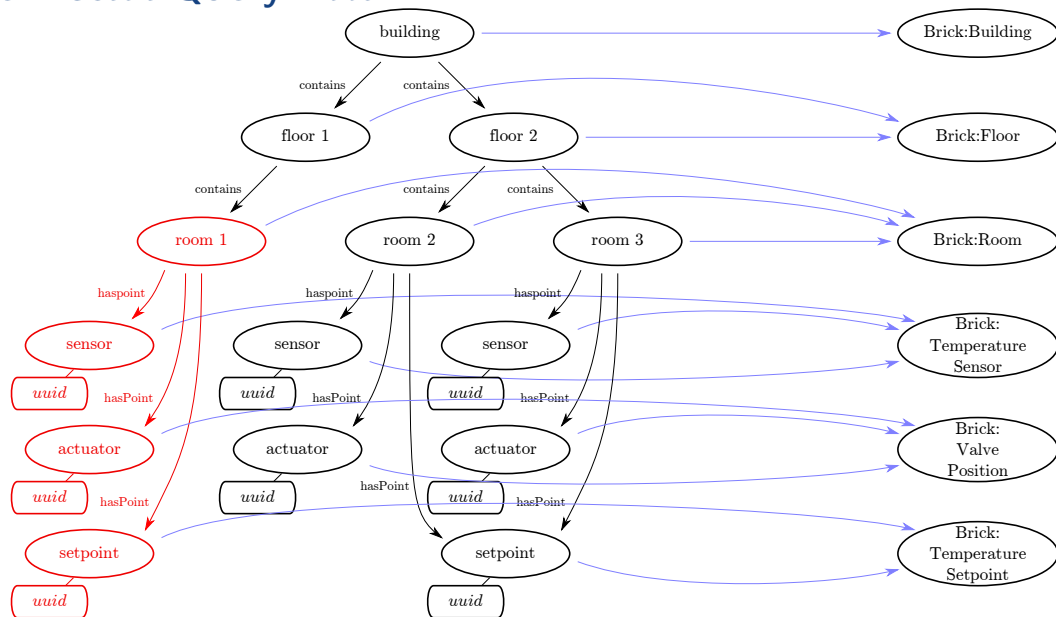
Demo: Thermostat Query

```
SELECT DISTINCT ?room_name ?sensor_uuid ?setpoint_uuid ?actuator_uuid
WHERE {
    ?room      rdf:type/brick:subClassOf* brick:Room .
    ?sensor    rdf:type/brick:subClassOf* brick:Temperature_Sensor .
    ?setpoint  rdf:type/brick:subClassOf* brick:Temperature_Setpoint .
    ?actuator  rdf:type/brick:subClassOf* brick:Radiator_Valve_Position .

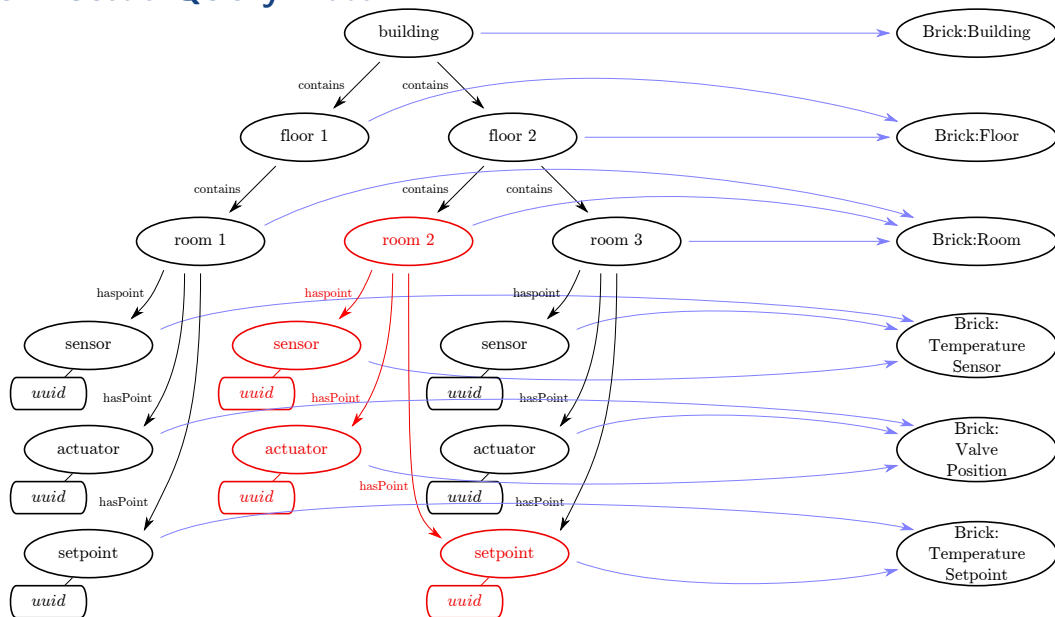
    ?sensor    brick:pointOf ?room .
    ?setpoint  brick:pointOf ?room .
    ?actuator  brick:pointOf ?room .

    ?room      brick:label ?room_name .
    ?sensor    brick:label ?sensor_uuid .
    ?setpoint  brick:label ?setpoint_uuid .
    ?actuator  brick:label ?actuator_uuid .
}
```

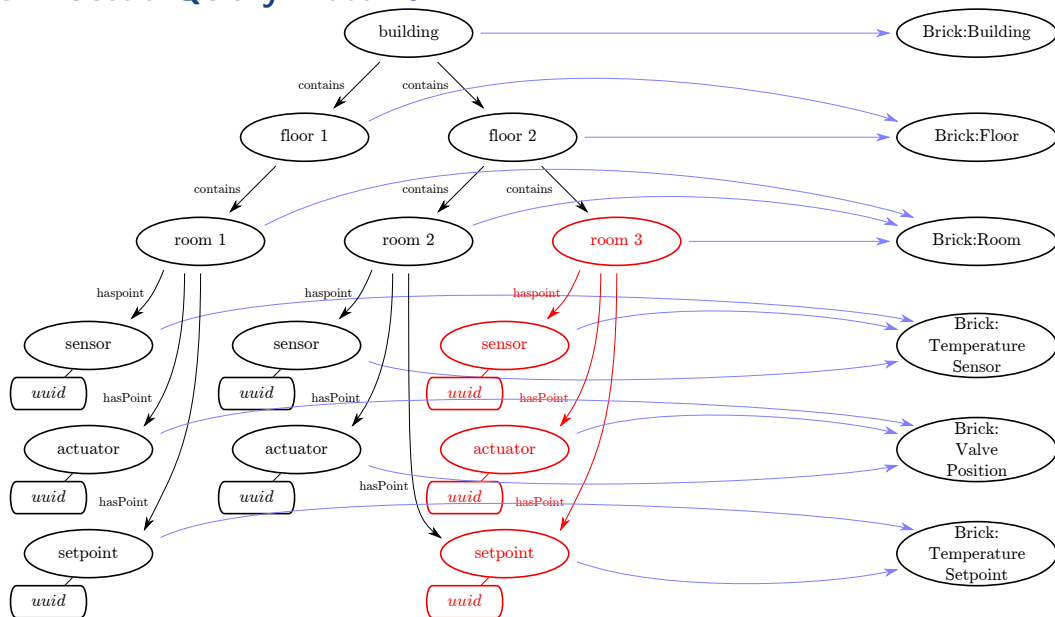
Thermostat Query Match 1



Thermostat Query Match 2



Thermostat Query Match 3



Thermostat Query Thermostat Result

```
[  
  [  
    "room 1",  
    "d0fcec33-af08-44a8-b74b-7a51f1902d13",  
    "c087d1ba-03f6-4c60-85de-b95d82e63248",  
    "5c2e2bdb-5142-464d-ad91-ae483633dfd6"  
  ], [  
    "room 2",  
    "18407632-8e81-4bb9-ae77-4ecc96f30f46",  
    "d2b54605-58b1-44ef-89f9-0a9adabbfb1",  
    "ec883cad-3235-46d5-a901-7bae169dda0a"  
  ], [  
    "room 3",  
    "e528f733-2e3a-4f65-93aa-6fb2e36d4b27",  
    "d2b54605-58b1-44ef-89f9-0a9adabbfb1",  
    "9be52aed-0e12-49b5-8680-47778b9b2adf"  
  ]  
]
```


Modification Deletion

```
DELETE {  
    ?actuator brick:label "ec883cad-3235-46d5-a901-7bae169dda0a" .  
}  
WHERE {  
    ?actuator rdf:type/brick:subClassOf* brick:Radiator_Valve_Position .  
    ?actuator brick:label "ec883cad-3235-46d5-a901-7bae169dda0a" .  
}
```

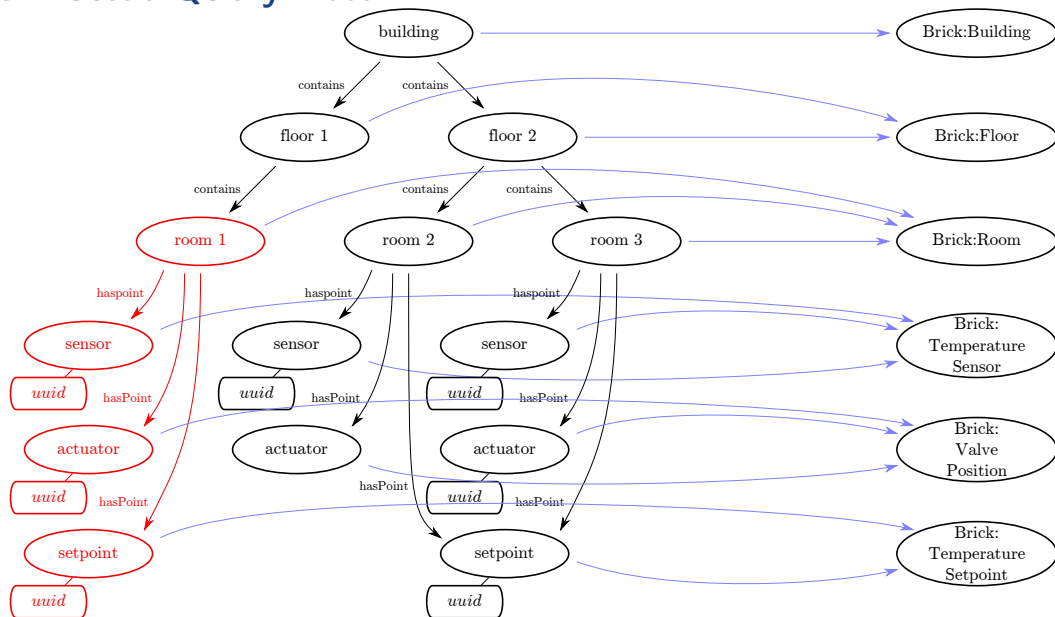
Deletion Pre Update



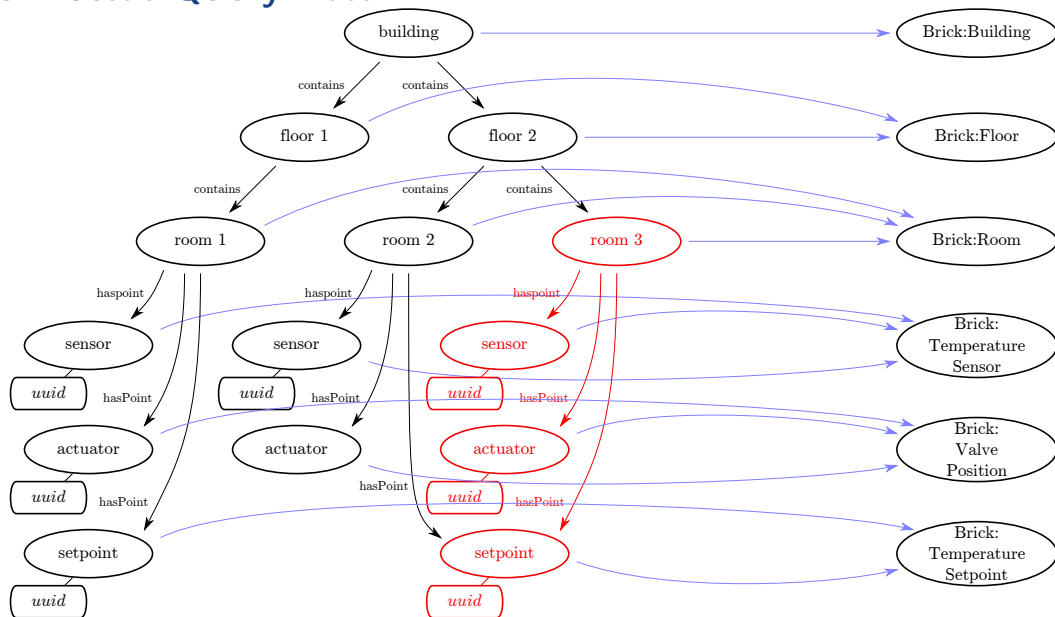
Deletion Post Update



Thermostat Query Match 1



Thermostat Query Match 2



Thermostat Query Results

```
[  
  [  
    "room 1",  
    "d0fcec33-af08-44a8-b74b-7a51f1902d13",  
    "c087d1ba-03f6-4c60-85de-b95d82e63248",  
    "5c2e2bdb-5142-464d-ad91-ae483633dfd6"  
  ], [  
    "room 3",  
    "e528f733-2e3a-4f65-93aa-6fb2e36d4b27",  
    "d2b54605-58b1-44ef-89f9-0a9adabbfb1",  
    "9be52aed-0e12-49b5-8680-47778b9b2adf"  
  ]  
]
```

Modification Inserts

Inserts can be done in the same way as you do deletes

Insert and delete clauses may be combined

Part 6:

Building a New Ontology

Ontology Construction Introduction

An ontology defines a the equivalent of a schema

A model contains instances of types and relates these instances to each other according to rules defined in such ontologies

A model may refer to multiple ontologies

Ontologies are simply RDF stores

They usually define a namespace and are stored in a file

Ontology Construction Defining Types

We can define a type by declaring the name a subclass of something else.

Everything that exists in the `owl:subClassOf` DAG rooted in `owl:Class` is a type.



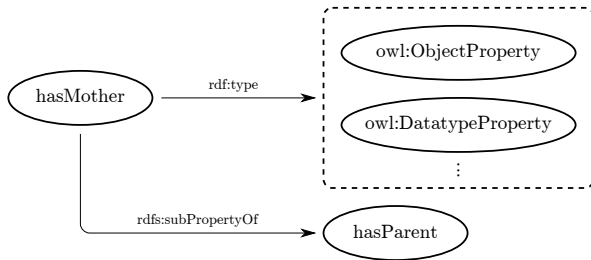
Ontology Construction

Defining Relationships

Relationships are defined through their use and their place in the DAG of types

Their use is based on whether it points at a literal (`owl:DatatypeProperty`) or not (`owl:ObjectProperty`)

Their place is based on the `rdfs:subPropertyOf` relation.




Note: In this example `hasMother` and `hasParent` are properties

Ontology Construction Relationship Restrictions

A relationship p can be restricted to only be allowed on some set S of subjects using `owl:domain` properties. Hereby you define the set of allowed triples of the form $(S,p,*)$

A relationship p can be restricted to only be allowed on some set O of objects using `owl:range` properties. Hereby you define the set of allowed triples of the form $(*,p,O)$




What happens when you combine them?

Ontology Construction Type Equivalence

We can state that two classes (e.g., types) are equivalent using `owl:equivalentType`

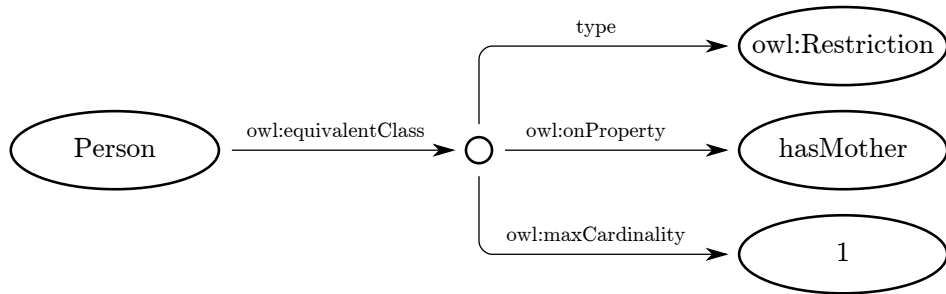
A common usecase is to bridge two ontologies in a model



Ontology Construction Restrictions Between Type and Relationship

We can limit how many triples are allowed to follow the pattern (s, p, *)

This is known as a restriction on cardinality



Note: In this example hasMother is a property

Ontology Construction Literals

The XSD namespace defines a set of types whose values can be encoded as literals

https://www.w3.org/2011/rdf-wg/wiki/XSD_Datatypes

Examples:

- ▶ int
- ▶ double
- ▶ boolean
- ▶ string
- ▶ dateTime

Part 7: Validation of Model

Role of a Validator

A validator tests the stated rules of the applied ontologies.

Examples:

- ▶ The `brick:isLocatedIn` relationship has as subject something that is a subclass of `brick:Equipment` and as object something that is a subclass of `brick:Location`
- ▶ Any instance of a `person:Human` has exactly two parents
- ▶ The `birthdate` property of an instance of a `person:Human` has the `date` type

Without a validator triples can be added at will

- ▶ Spelling mistakes are accepted
- ▶ Entities can be connected in any way
(in general, the order of definitions don't matter)
- ▶ Entities can be floating

Part 8: Reasoners

Role of a Reasoner

A reasoner can apply rules

Examples:

- ▶ Given the facts that **Marge is a Female** and **Marge is parent of Maggie** a reasoner can conclude that **Marge is the mother of Maggie**
- ▶ Given a fact stating that **some entity is a temperature sensor** and another stating that **temperature sensor is a subclass of sensor** a reasoner can conclude that **this entity is a sensor**

Note: Each of these require a rule

If you define such a rule, a reasoner will be able to apply it

Consequences of using Reasoner

When to apply the reasoner:

1. When inserting
2. On triplet store before querying
3. While querying

Instead of:

```
?maggie person:hasName "Maggie" .  
?mother person:isParentOf ?maggie .  
?mother person:isA person:Female .
```

We can write:

```
?maggie person:hasName "Maggie" .  
?mother person:isMotherOf ?maggie .
```

Consequences of using Reasoner

When to apply the reasoner:

1. When inserting
2. On triplet store before querying
3. While querying

Instead of:

```
?sensor rdf:type/rdf:subClassOf* brick:Sensor .
```

We can write:

```
?sensor rdf:type brick:Sensor .
```

Part 9: Model Management

Hosting a Model

Many options:

1. Filesystem
2. Webserver
3. Dedicated RDF server
 - ▶ HodDB – Fast, read-only, stability issues
 - ▶ Fuseki – Supports writes

At SDU we are working on

- ▶ rdfserv – slow, supports writes and multiple protocols
- ▶ A next-generation RDF server supporting query-based subscriptions

Interacting with a Model

Most popular high-level languages have libraries for manipulating and querying a triple store

- ▶ ***Python:*** `rdflib`
- ▶ ...

Typical qualities:

- ▶ No validation
- ▶ Limited or no reasoners
- ▶ "Slow" query resolution available
 - ▶ Simple queries are fast
 - ▶ Large complex queries takes forever

This is where dedicated RDF servers come into play!

Further Reading

RDF and OWL:

- ▶ Google has plenty of information

Brick:

- ▶ Official homepage at <http://brickschema.org>

Evaluating SparQL Queries:

- ▶ **Library** For python there is the rdflib module at <https://github.com/RDFLib/rdflib>
- ▶ **Web Service**
 - ▶ Fuseki: https://jena.apache.org/documentation/serving_data/
 - ▶ HodDB: <https://hoddb.org>

Questions?

